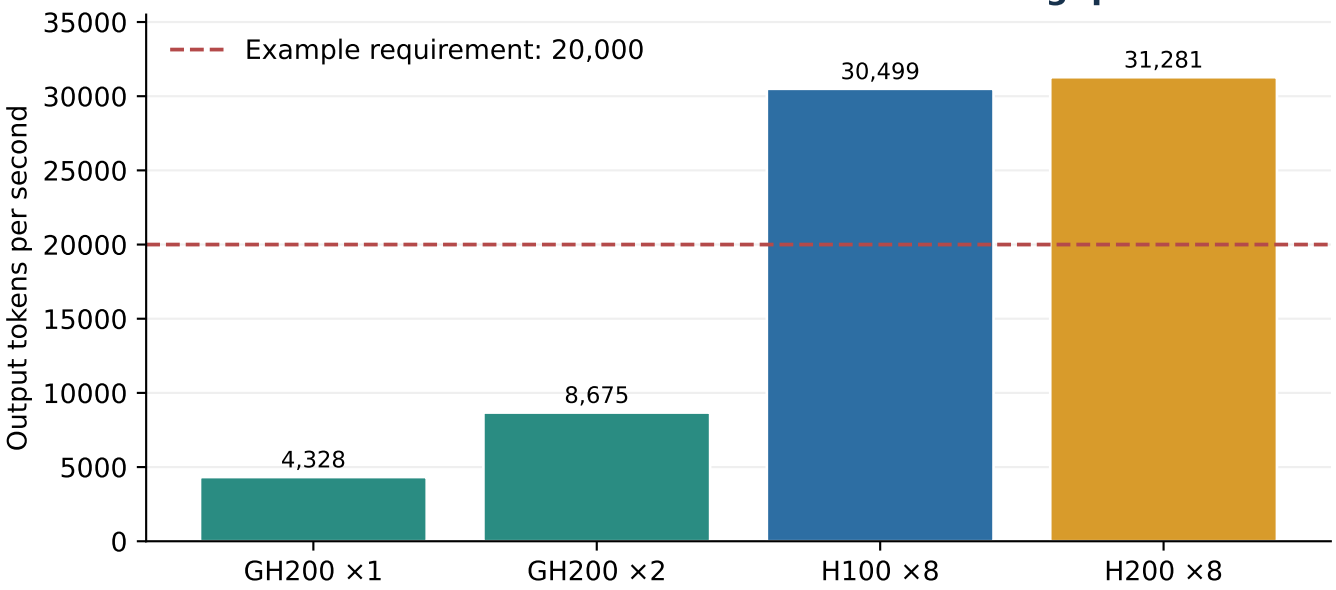
S2: measured Llama 2 70B Server throughput
Output tokens per second
35000
30000
25000
20000
15000
10000
5000
0
Example requirement: 20,000
4,328
8,675
30,499
31,281
GH200 ×1
GH200 ×2
H100 ×8
H200 ×8